

# Performance Metrics- Classification



**FAIRLEIGH  
DICKINSON  
UNIVERSITY**

# ML Flow

**Given:** labeled training data  $X, Y = \{< x_i, y_i >\}_{i=1}^n$

- Assumes each  $x_i \sim D(X)$  with  $y_i = f_{target}(x_i)$

**Train the model:**

```
model ← classifier.train(X, Y )
```

**Apply the model to new data:**

- Given: new unlabeled instance  $x \sim D(X)$

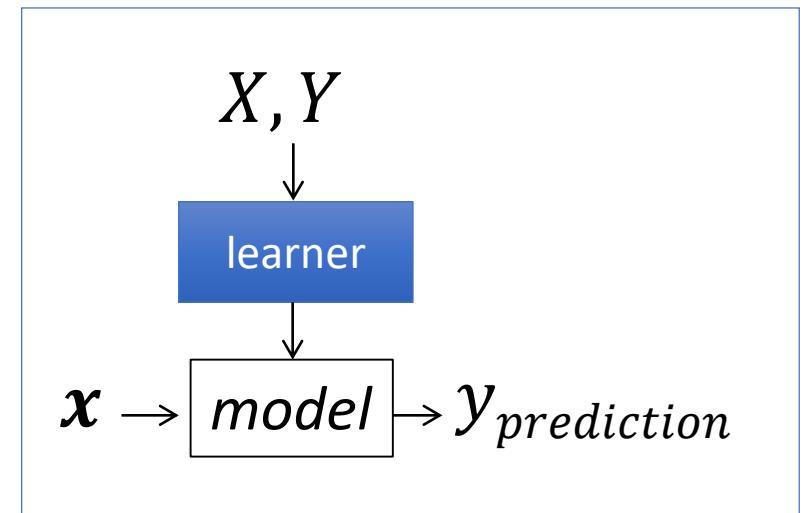
```
y_prediction ← model.predict(x)
```

**Key questions:**

How to determine the quality of the model?

(i) measuring performance

(ii) understanding the significance of the results (is it better than other models?)



**FAIRLEIGH  
DICKINSON  
UNIVERSITY**



# Why performance metrics are important?

- Training objective (cost function) is only a proxy for real world objectives.
- Metrics help capture a business goal into a quantitative target (not all errors are equal).
- Helps organize ML team effort towards that target.
  - Generally in the form of improving that metric on the dev set.
- Useful to quantify the “gap” between:
  - Desired performance and baseline (estimate effort initially).
  - Desired performance and current performance.
  - Measure progress over time.
- Useful for lower level tasks and debugging (e.g. diagnosing bias vs variance).
- Ideally training objective should be the metric, but not always possible. Still, metrics are useful and important for evaluation.

# Metrics for Performance Evaluation

- Focus on the predictive capability of a model
  - Rather than how fast it takes to classify or build models, scalability, etc.
- **Confusion Matrix:**

		PREDICTED CLASS	
		Class=Yes	Class>No
ACTUAL CLASS	Class=Yes	a: TP	b: FN
	Class>No	c: FP	d: TN

a: TP (true positive)

b: FN (false negative)

c: FP (false positive)

d: TN (true negative)



# Metrics for Performance Evaluation

		PREDICTED CLASS	
ACTUAL CLASS		Class=Yes	Class>No
	Class=Yes	a: TP	b: FN
	Class>No	c: FP	d: TN

- a: TP (true positive)
- b: FN (false negative)
- c: FP (false positive)
- d: TN (true negative)

**Most widely-used metric**

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$

Error=1- Accuracy



**FAIRLEIGH  
DICKINSON  
UNIVERSITY**

# Binary Classification

- $x$  is input
- $y$  is binary output (0/1)
- Model is  $\hat{y} = h(x)$
- Two types of models
  - Models that output a categorical class directly (K-nearest neighbor, Decision tree)
  - Models that output a real valued score (Logistic Regression, NN, SVM)
    - Score could be probability (LR, NN), margin:distance from the decision boundary(SVM)
    - Need to pick a threshold
    - We focus on this type (the other type can be interpreted as an instance)



# Score Based Models

Score = 1



●	Positive example
○	Negative example

Example of Score: Output of logistic regression.  
For most metrics: Only ranking matters.  
If too many examples: Plot class-wise histogram.

$$\text{Prevalence} = \frac{\# \text{ positive samples}}{\# \text{ positive} + \# \text{ negatives samples}}$$

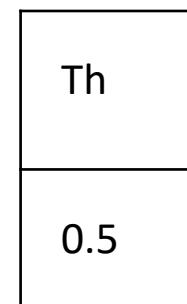
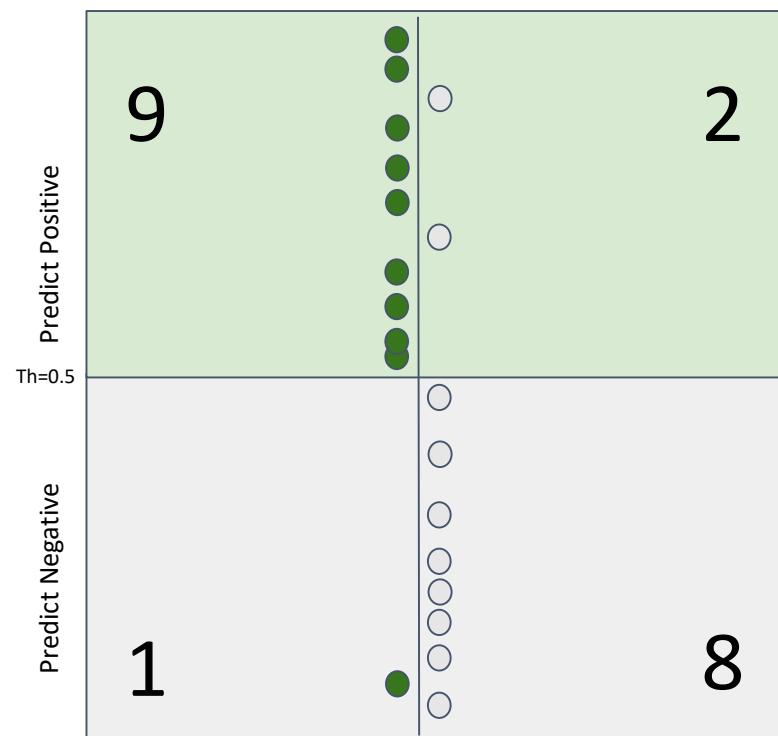
Prevalence: tells you **how imbalanced** your dataset is



FAIRLEIGH  
DICKINSON  
UNIVERSITY

# Confusion Matrix

Label Positive      Label Negative



## Properties:

- Total sum is fixed (population).
- Column sums are fixed (class-wise population).
- Quality of model & threshold decide how columns are split into rows.
- We want diagonals to be “heavy”, off diagonals to be “light”.

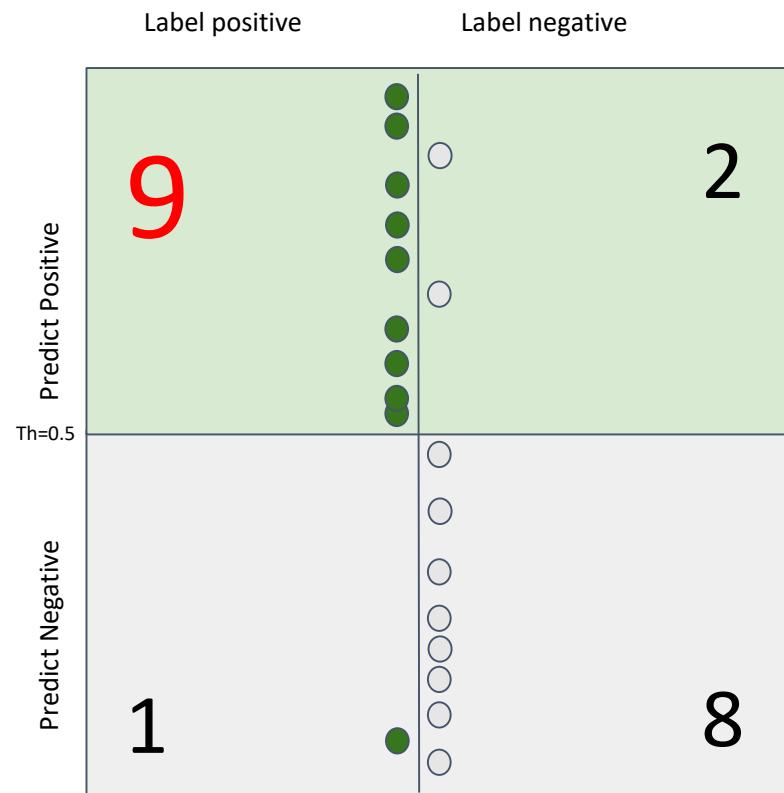


FAIRLEIGH  
DICKINSON  
UNIVERSITY

# True Positives



FAIRLEIGH  
DICKINSON  
U N I V E R S I T Y

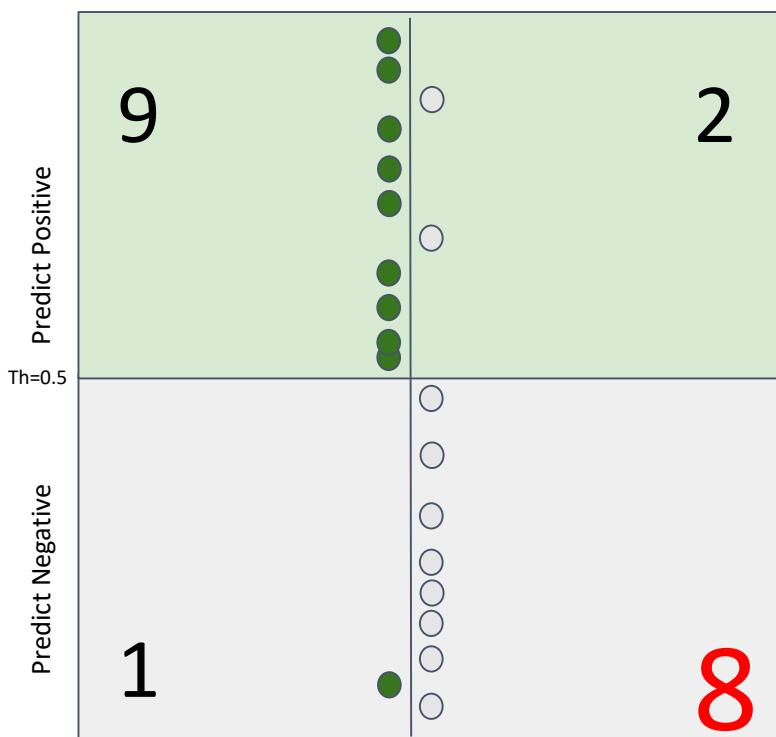


Th	TP
0.5	9

# True Negatives

Label positive

Label negative



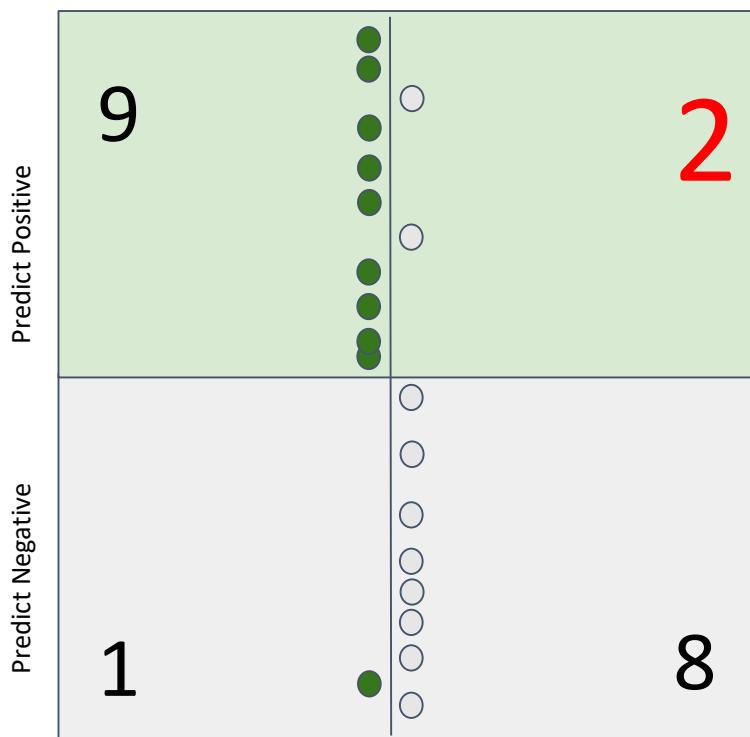
Th	TP	TN
0.5	9	8



FAIRLEIGH  
DICKINSON  
UNIVERSITY

# False Positives

Label positive      Label negative



Th	TP	TN	FP
0.5	9	8	2



FAIRLEIGH  
DICKINSON  
U N I V E R S I T Y

# False Negatives

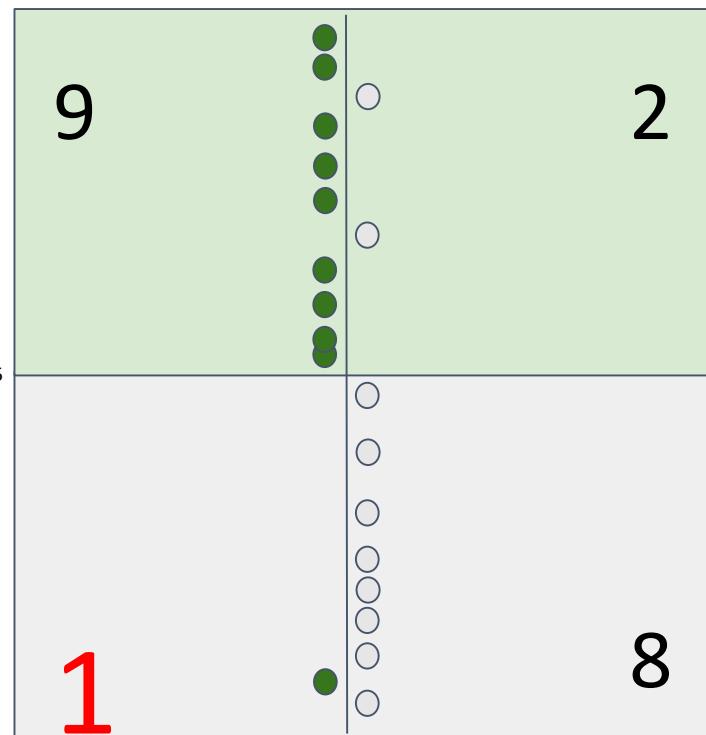
Label positive

Label negative

Predict Positive

Th=0.5

Predict Negative



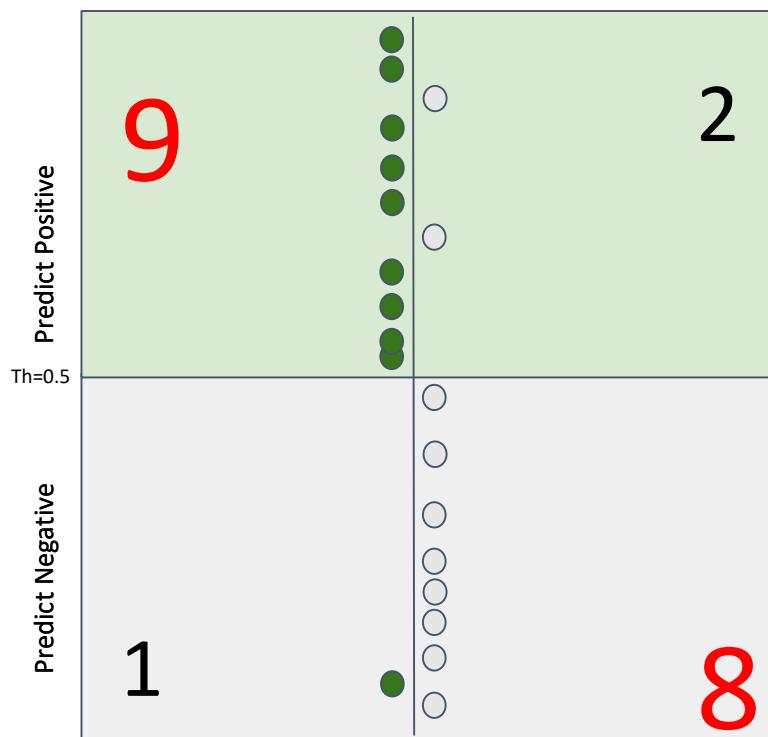
Th	TP	TN	FP	FN
0.5	9	8	2	1



FAIRLEIGH  
DICKINSON  
UNIVERSITY

# Accuracy

Label positive      Label negative



Th	TP	TN	FP	FN	Acc
0.5	9	8	2	1	.85



FAIRLEIGH  
DICKINSON  
UNIVERSITY

Equivalent to 0-1 Loss!

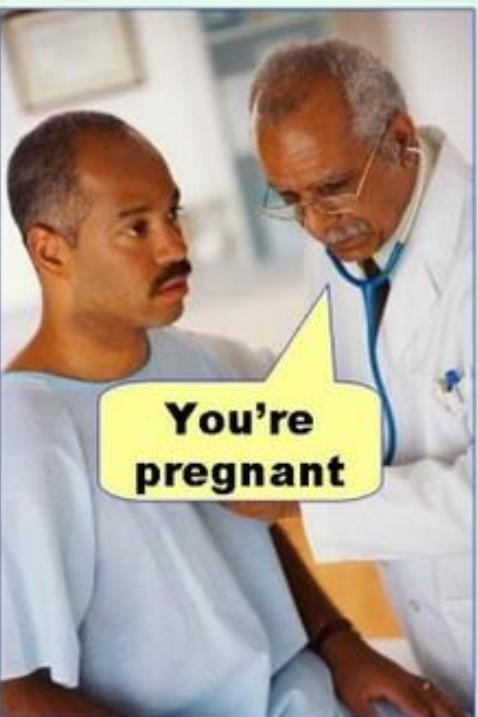
Any possible problem with it?

# FP and FN also called Type-1 and Type-2 errors



FAIRLEIGH  
DICKINSON  
UNIVERSITY

**Type I error**  
(false positive)



**Type II error**  
(false negative)



# Limitation of Accuracy



FAIRLEIGH  
DICKINSON  
UNIVERSITY

- Consider a 2-class problem (binary)
  - Number of Class 0 examples = 9990
  - Number of Class 1 examples = 10
- If model predicts everything to be class 0, accuracy is  $9990/10000 = 99.9\%$ 
  - Accuracy is misleading because model does not detect any class 1 example

# Alternate Metrics

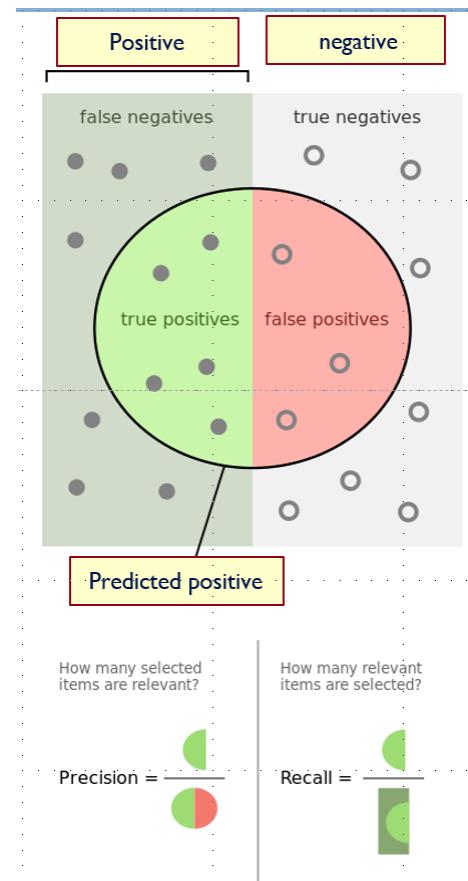
- If the Binary classification problem is biased
  - In many problems most examples are negative
- Or, in multiclass classification
  - The distribution over labels is often non-uniform
- Simple accuracy is not a useful metric.
  - Often we resort to task specific metrics
- We require to involve **Recall** and **Precision**

**Recall:** # true positives/ # all positive

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

**Precision:** # true positives/ # predicted positive

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$



# Alternate Metrics(cont.)

- Given a dataset of P positive instances and N negative instances:

		Predicted Class	
		Yes	No
Actual Class	Yes	TP	FN
	No	FP	TN

The notion of a confusion matrix can be usefully extended to the multiclass case. (i,j) cell indicate how many of the i-labeled examples were predicted to be j.

- Imagine using classifier to identify positive cases (i.e., for information retrieval)

$$\text{precision} = \frac{TP}{TP + FP}$$

Probability that a randomly selected positive prediction is indeed positive

$$\text{accuracy} = \frac{TP + TN}{P + N}$$

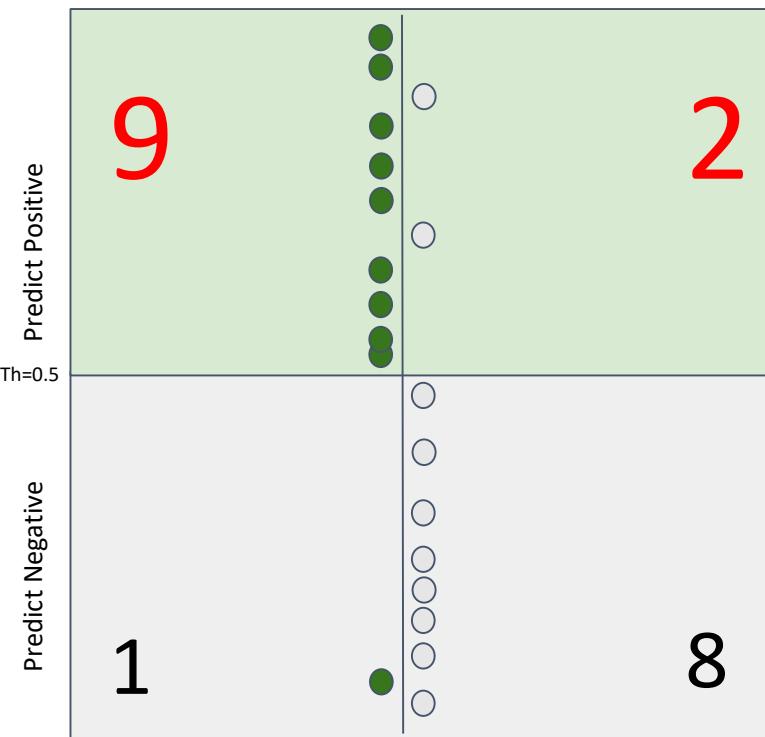
$$\text{recall} = \frac{TP}{TP + FN}$$

Probability that a randomly selected positive is identified



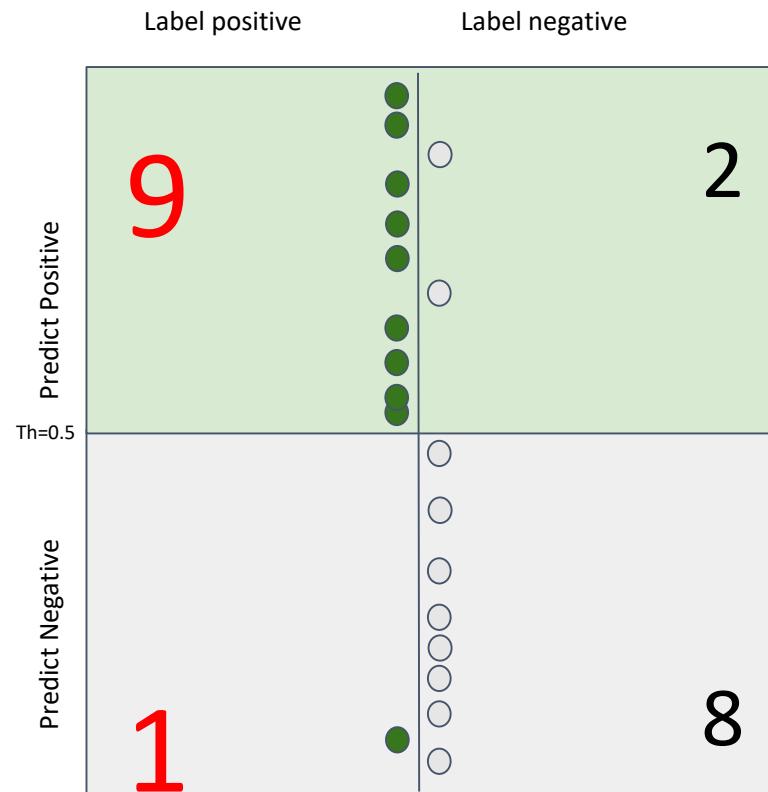
# Precision

Label positive      Label negative



Th	TP	TN	FP	FN	Acc	Pr
0.5	9	8	2	1	.85	.81

# Recall (Sensitivity)



Th	TP	TN	FP	FN	Acc	Recall
0.5	9	8	2	1	.85	.9

Trivial 100% recall = pull everybody above the threshold.

Trivial 100% precision = push everybody below the threshold except 1 green on top.  
(Hopefully no gray above it!)

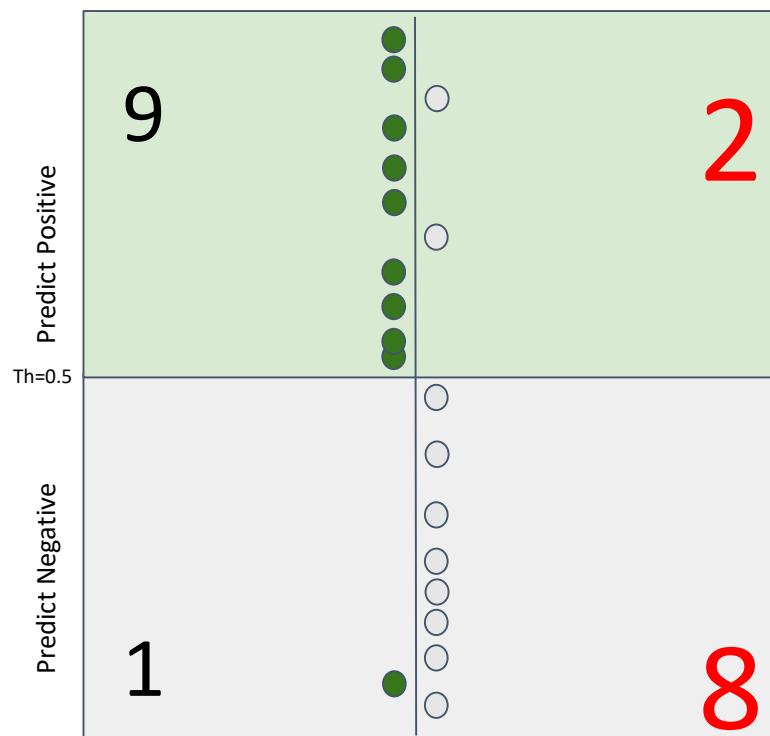
Striving for good precision with 100% recall = pulling up the lowest green as high as possible in the ranking.  
Striving for good recall with 100% precision = pushing down the top gray as low as possible in the ranking.



# Negative Recall (Specificity)

Label positive

Label negative



Th	TP	TN	FP	FN	Acc	Pr	Recall	Spec
0.5	9	8	2	1	.85	.81	.9	0.8



# Examples- Imbalanced Data

Scenario	What's more important	Reason
Medical diagnosis	Recall	Missing a positive case is dangerous
Fraud detection	Both	Missing fraud = loss; false alarms = frustration
Spam filtering	Precision	Avoid marking real emails as spam
Cybersecurity	Recall (with balanced Precision)	Missing an attack can be critical
Search engines	Precision	Users prefer fewer but relevant results



# Example: Medical Diagnosis(e.g. Cancer Detection)

- **Goal:** Identify all patients who have cancer.
- **Challenge:** Very few positive cases → class imbalance.
- **Why Recall matters most:**
  - Missing a patient with cancer (**false negative**) can be life-threatening.
  - It's better to raise some false alarms (low precision) than to miss actual positives.

## Example:

1000 patients → 10 have cancer

- Model predicts 8 of them correctly (Recall = 0.8), but also wrongly flags 20 healthy ones (Precision = 0.29).
- Accuracy might look 98%, but Recall tells us how well we're catching the real cases.



# Example: Spam Email Filtering

- **Goal:** Identify spam emails automatically.
- **Why Precision matters most:**
  - If Precision is low, important legitimate emails get marked as spam (**false positives**), which frustrates users.
  - Recall can be slightly lower (missing a few spam emails) — that's acceptable.
  - Gmail's filters aim for very high precision so that real emails are never lost, even if a few spam ones slip through.



# Example: Credit Card Fraud Detection

- **Goal:** Detect fraudulent transactions.
- **Challenge:** Fraud cases are <1% of total transactions.
- **Why both Precision and Recall matter:**
  - Low Recall: Miss fraud → financial loss.
  - Low Precision: Too many false alarms → annoy customers and waste resources.
- So, the goal is to find a good balance (often via **F1-score**) that minimizes both risks.



# F1 Score

- Precision and recall often **conflict**
  - Increasing **recall** (catching more positives) may reduce precision (more false alarms).
  - Increasing **precision** (fewer false alarms) may reduce recall (missed positives).
- **Accuracy** is misleading on imbalanced data.
  - It makes sense to consider Recall and Precision together or combine them into a single metric.
- **F1** provides a *single, interpretable metric* when both Precision and Recall are important.
  - It **penalizes extreme imbalance** — a high F1 requires both precision and recall to be reasonably good.
- **F1 Score** is a measure that **combines precision and recall** is the harmonic mean of precision and recall.

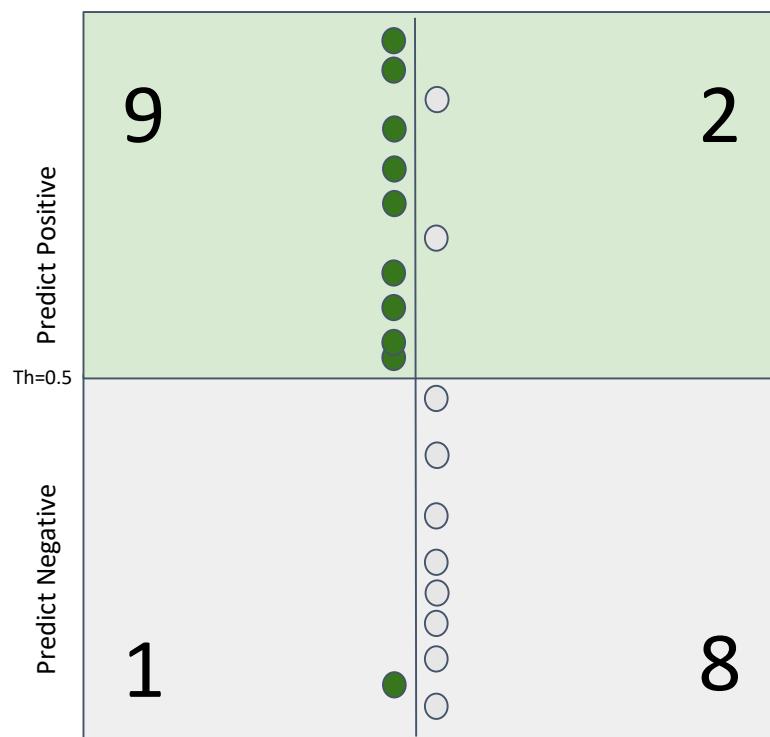
$$F_1 = \left( \frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} \right) = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



# F1 Score

Label positive

Label negative



Th	TP	TN	FP	FN	Acc	Pr	Recall	Spec	F1
0.5	9	8	2	1	.85	.81	.9	.8	.857

$$F_1 = \left( \frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} \right) = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}.$$



# Example

Suppose we have a medical test for detecting a rare disease.

Metric	Value	Interpretation
Recall	0.5	50% of all actual patients were detected
Precision	0.8	80% of detected cases were correct
F1 Score	?	Combines the two

$$F1 = 2 \times (0.5 * 0.8) / (0.5 + 0.8) = 0.615$$

**F1 = 0.615**, showing moderate performance — better than either metric alone.



# Why harmonic mean (not arithmetic)?

Because the harmonic mean **penalizes imbalance** — if one of precision or recall is very low, F1 drops sharply.

Example:

Precision = 1.0, Recall = 0.1 → F1 = 0.18

Precision = 0.55, Recall = 0.55 → F1 = 0.55

So, the F1 score emphasizes **consistency** between precision and recall.



# How to improve F1?

Approach	Description
<b>Threshold tuning</b>	Adjust classification threshold (e.g., $0.5 \rightarrow 0.3$ ) to increase recall or precision as needed
<b>Class weighting</b>	Penalize false negatives or positives differently during training
<b>Resampling</b>	Oversample minority class or undersample majority class
<b>Better features or model</b>	Use richer data, embedding, or ensemble models to improve overall discrimination

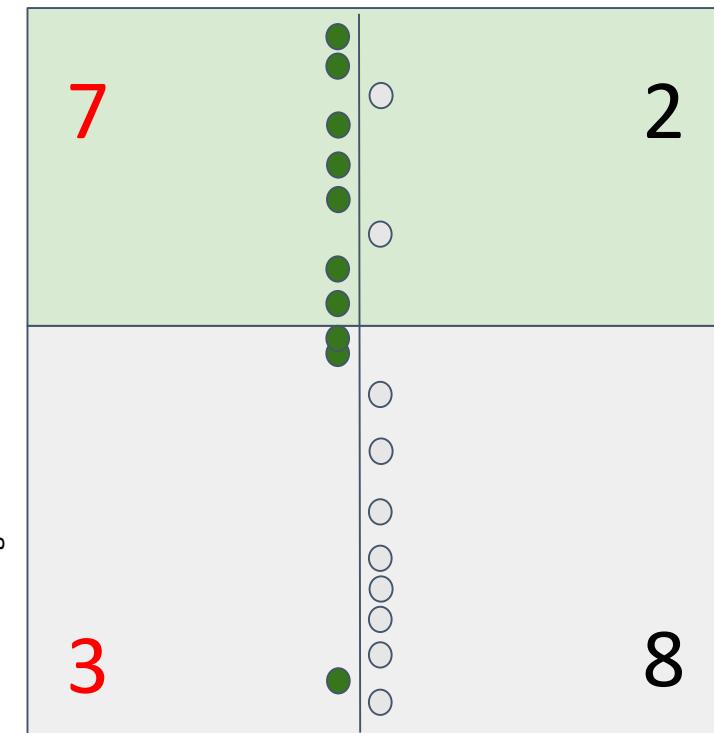


# Changing threshold

Label positive

Label negative

Predict Positive



Th	TP	TN	FP	FN	Acc	Pr	Recall	Spec	F1
0.6	7	8	2	3	.75	.77	.7	.8	.733

# effective thresholds = # examples + 1

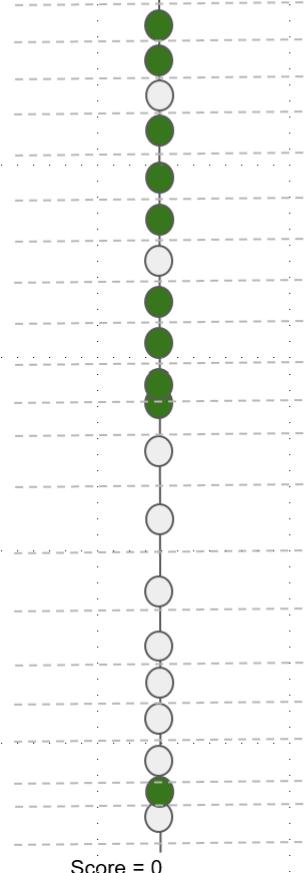


# Threshold Tuning

Threshold Scanning

Score = 1

Threshold = 1.00



Threshold = 0.00

Score = 0

Threshold	TP	TN	FP	FN	Accuracy	Precision	Recall	Specificity	F1
1.00	0	10	0	10	0.50	1	0	1	0
0.95	1	10	0	9	0.55	1	0.1	1	0.182
0.90	2	10	0	8	0.60	1	0.2	1	0.333
0.85	2	9	1	8	0.55	0.667	0.2	0.9	0.308
0.80	3	9	1	7	0.60	0.750	0.3	0.9	0.429
0.75	4	9	1	6	0.65	0.800	0.4	0.9	0.533
0.70	5	9	1	5	0.70	0.833	0.5	0.9	0.625
0.65	5	8	2	5	0.65	0.714	0.5	0.8	0.588
0.60	6	8	2	4	0.70	0.750	0.6	0.8	0.667
0.55	7	8	2	3	0.75	0.778	0.7	0.8	0.737
0.50	8	8	2	2	0.80	0.800	0.8	0.8	0.800
0.45	9	8	2	1	0.85	0.818	0.9	0.8	0.857
0.40	9	7	3	1	0.80	0.750	0.9	0.7	0.818
0.35	9	6	4	1	0.75	0.692	0.9	0.6	0.783
0.30	9	5	5	1	0.70	0.643	0.9	0.5	0.750
0.25	9	4	6	1	0.65	0.600	0.9	0.4	0.720
0.20	9	3	7	1	0.60	0.562	0.9	0.3	0.692
0.15	9	2	8	1	0.55	0.529	0.9	0.2	0.667
0.10	9	1	9	1	0.50	0.500	0.9	0.1	0.643
0.05	10	1	9	0	0.55	0.526	1	0.1	0.690
0.00	10	0	10	0	0.50	0.500	1	0	0.667

# Question?



**FAIRLEIGH  
DICKINSON  
UNIVERSITY**